

DigitalOcean hosted inference

Interim evidence rebuild - 28 August 2026

EARLIER REPORT WITHDRAWN

Do not use prior PDFs or plots for decisions

NOT COMPLETE

NOT A PRODUCTION QUALIFICATION

19/44	3/44	100/176
adaptive-load cells with a repeat-confirmed rate	120-second fixed-rate tests that passed	endpoint-capability checks with complete evidence

The fixed-rate result is 3 passes, 38 measured failures, and 3 tests that could not establish a reliable baseline. Finishing a request schedule does not count as a pass.

Six-hour time-of-day variation has not been measured yet. This edition reserves a section for that matched panel and makes no 24-hour, daily, or diurnal claim.

Scope: 11 exact DigitalOcean-hosted open-model endpoints. Commercial pass-through routes are excluded. Capacity and fixed-rate source identifiers are printed in the provenance section.

How to read the experiments

Adaptive load search raises offered traffic while the endpoint is healthy, reduces traffic after degradation, and requires three separated healthy confirmations before publishing a numeric rate. A confirmed lower bound is not a maximum.

The 120-second fixed-rate stability test holds one rate for four adjacent 30-second blocks, then checks recovery. A pass requires every registered reliability, latency, queueing, usage, quality, and recovery condition to pass. Its four-block intervals are exploratory because serial correlation is not modeled.

The capacity sources used different recipes. The 2026-08-24 source used a 32K long prompt; the 2026-08-28 source used a 100K long prompt (50K for Minimax). They remain separate in every table and chart.

Source	Workload	Exact registered recipe
do-capacity-20260828-r2	short prompt / short answer	256-token prompt -> 128-token answer target
do-capacity-20260828-r2	100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5)
do-capacity-20260828-r2	short prompt / long answer	256-token prompt -> 4,096-token answer target
do-capacity-20260828-r2	seeded multi-workload mix	seeded four-way mix: 256->128, 100K/50K->128, 256->4,096, and 1,024->512 JSON
do-sixhour-aimd-20260824-r1	short prompt / short answer	short exact-answer prompt -> 64-token answer ceiling
do-sixhour-aimd-20260824-r1	32K-token prompt / short answer	32,000-token prompt -> 64-token answer ceiling
do-sixhour-aimd-20260824-r1	short prompt / long answer	short prompt -> 1,024-word target / 2,048-token ceiling
do-sixhour-aimd-20260824-r1	seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling

The combined capacity table uses 21 exact endpoint-by-workload controllers from the corrected 2026-08-28 closure and 23 matching cells from the earlier verified six-hour campaign. The corrected closure's four-hour guard censored those fallback cells before start; none is presented as new evidence.

Endpoint evidence map

Counts below are evidence counts, not a provider ranking. A production decision must match the application's exact workload and required capability.

Exact endpoint	Documented context	Repeat-confirmed load rates	Fixed-rate stability passes	Capability checks passed
deepseek-v4-flash-0731	1048576	4/4	1/4	12/18
gemma-4-31B-it	256000	3/4	1/4	13/18
glm-5.2	262144	1/4	0/4	8/25
kimi-k3	not documented	1/4	0/4	2/18
mimo-v2.5-pro	262144	2/4	0/4	10/25
minimax-m2.5	65536	3/4	0/4	1/18
nemotron-3-ultra-550b	131072	2/4	0/4	9/18
nvidia-nemotron-3-super-120b	1000000	0/4	0/4	7/18
openai-gpt-oss-120b	128000	2/4	0/4	6/18
qwen3.5-397b-a17b	131072	0/4	0/4	1/18
qwen3.8-max	1000000	1/4	1/4	18/25

Separate static verification

Campaign do-complete-20260827-r1 is a separate, hash-bound verification layer (0b2a4772c5c444d25d66be7f851488456d09551c4ce5fd6575c4017d793796ac). It contributes static, caching, capability, context, output, quality, and latency evidence only. It contributes no adaptive-load or sustained-capacity claim. Cells not reached before its wall-time guard remain explicitly labelled.


Exact endpoint	Suites reached	Attempts	Successes	Observed state
deepseek-v4-flash-0731	cache, capability, context, output, warmup	84	76	measured
gemma-4-31B-it	latency, warmup	50	49	measured
glm-5.2	cache, interactions, quality, warmup	65	65	measured
kimi-k3	interactions, latency, quality, warmup	152	19	measured
mimo-v2.5-pro	context, quality, warmup	39	37	measured
minimax-m2.5	cache, warmup	25	25	measured
nemotron-3-ultra-550b	interactions, latency, output, warmup	73	71	measured
nvidia-nemotron-3-super-120b	cache, context, output, quality, warmup	83	59	measured
openai-gpt-oss-120b	warmup	1	1	measured
qwen3.5-397b-a17b	cache, latency, warmup	71	63	measured
qwen3.8-max	warmup	1	1	measured

Automatic exact-prefix cache verification

DigitalOcean-hosted open models apply best-effort exact-prefix caching automatically. The matched pairs below report observed cached-token counters, TTFT when the transport exposed it, and settled-cost ratios. A missing row means the verification campaign did not reach that endpoint; it is not a claim that caching is unavailable.

Endpoint	Pairs	Hits / misses	Cached tokens	TTFT cached / uncached	Cost ratio
deepseek-v4-flash-0731	12/12	3/9	24576	1.08 / 0.996 s	0.83x
glm-5.2	12/12	11/1	89408	not reported	0.24x
minimax-m2.5	12/12	12/0	82448	not reported	0.37x
nvidia-nemotron-3-super-120b	12/12	11/1	88704	not reported	0.29x
qwen3.5-397b-a17b	12/12	11/1	87120	not reported	0.31x


32K-era adaptive load search (2026-08-24 source)



■ not sourced from this campaign
 ■ ran; no numeric bound
 ■ no healthy result
 ■ observed once; not confirmed
 ■ three separated confirmations


What was run	Five-second open-loop epochs, then three separated confirmations where possible. Exact recipes: short exact answer -> 64-token ceiling; 32,000-token prompt -> 64-token ceiling; short prompt -> 1,024-word target / 2,048-token ceiling; seeded five-way exact-answer, 4,096-context, long-answer, JSON, and tool-call mix.
What it proves	Green cells establish a repeatedly passing lower bound for that exact endpoint and recipe. Other colors are explicit non-results or weaker evidence.
What it does not prove	No chart on this page establishes a provider-wide maximum, production approval, 24-hour behavior, or diurnal behavior.

100K-era adaptive load search (2026-08-28 source)




What was run	Five-second open-loop epochs, then three separated confirmations where possible. Exact recipes: 256 -> 128 tokens; 100,000 -> 128 tokens (50,000 for Minimax); 256 -> 4,096 tokens; seeded four-way short, 100K/50K, long-answer, and JSON mix.
What it proves	Only green cells establish a repeat-confirmed lower bound. These results are separate from the 32K source and cannot be substituted for it.
What it does not prove	No chart on this page establishes a provider-wide maximum, production approval, 24-hour behavior, or diurnal behavior.

120-second fixed-rate stability results




passed every acceptance check test ran; one or more checks failed reliable baseline could not be established

What was run	One printed rate per endpoint and historical 32K-era workload for four adjacent 30-second blocks, plus recovery. The four recipes match the 32K chart above.
What it proves	PASS means every registered check passed at that exact rate for 120 seconds. It does not prove six-hour stability or a higher rate.
What it does not prove	No chart on this page establishes a provider-wide maximum, production approval, 24-hour behavior, or diurnal behavior.



What was counted	Every registered reason attached to a failed or could-not-start fixed-rate cell. One cell can contribute to several bars, so the bars do not sum to 41.
How to use it	The bars show which engineering failure modes dominated; they are not independent events and do not sum to the number of failed cells.
What it does not prove	No chart on this page establishes a provider-wide maximum, production approval, 24-hour behavior, or diurnal behavior.

DigitalOcean functional capability evidence



What was tested	Valid and malformed capability probes were scored independently where evidence existed.
What it proves	Every cell carries an explicit evidence state: passed, degraded, failed, documented unavailable, or not measured.
What it does not prove	No chart on this page establishes a provider-wide maximum, production approval, 24-hour behavior, or diurnal behavior.

Six-hour matched variation panel - pending

EVIDENCE STATUS

PENDING - no verified matched six-hour panel

No six-hour time-of-day panel exists in the retained evidence. The closure experiment must repeat the same endpoint, workload recipe, offered rate, region, and acceptance checks at predeclared times across a six-hour window.

Matched timeline

Hour 0	Registered checkpoints	Hour 6
Lock endpoint + recipe + rate + region; record the low-load reference.	Repeat the identical matched panel only at times predeclared in the closure plan; preserve request-level receipts.	Run the final matched panel and recovery check; close the evidence set.

Measurements required at every panel

Reliability	success rate; rate-limit rate; timeout plus server-error rate
Responsiveness	time to first token; end-to-end latency; achieved request rate; queue growth
Efficiency	completion throughput; prompt and completion usage completeness; cache counters where exposed
Quality and recovery	matched quality result during load; latency and quality after load

Claims boundary

After verification, this panel can support only a six-hour within-run variation statement. It cannot support a 24-hour, daily, diurnal, provider-wide maximum, or production-approval claim.

deepseek-v4-flash-0731

Model / API	deepseek-v4-flash-0731 / chat_completions
Region / API version	not reported / v1
Context / max output	1048576 / 1048576 tokens
Input / output price	\$0.080 / \$0.252 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	short exact-answer prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	repeatedly passed through at least 32 RPS; ceiling not found	PASS at 1 RPS
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	repeatedly passed through at least 0.5 RPS; ceiling not found	not run for the 100K recipe
short prompt / long answer	256-token prompt -> 4,096-token answer target do-capacity-20260828-r2	repeatedly passed through at least 1 RPS; ceiling not found	failed at 1.6 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	repeatedly passed at 0.0666667 RPS; degraded by 0.2 RPS	failed at 1.6 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed transport degraded	not scored
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 2722; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

gemma-4-31B-it

Model / API	gemma-4-31B-it / chat_completions
Region / API version	not reported / v1
Context / max output	256000 / 8192 tokens
Input / output price	\$0.18 / \$0.50 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	256-token prompt -> 128-token answer target do-capacity-20260828-r2	repeatedly passed at 0.25 RPS; degraded by 2.5 RPS	PASS at 1 RPS
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	not run for the 100K recipe
short prompt / long answer	256-token prompt -> 4,096-token answer target do-capacity-20260828-r2	repeatedly passed through at least 1 RPS; ceiling not found	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
seeded multi-workload mix	seeded four-way mix: 256->128, 100K/50K->128, 256->4,096, and 1,024->512 JSON do-capacity-20260828-r2	repeatedly passed through at least 2 RPS; ceiling not found	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed supported	passed
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 251873; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed functional; tool calling: observed functional; structured output: observed functional

glm-5.2

Model / API	glm-5.2 / chat_completions
Region / API version	not reported / v1
Context / max output	262144 / 262144 tokens
Input / output price	\$0.70 / \$2.20 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	256-token prompt -> 128-token answer target do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	failed at 1 RPS: p95 end-to-end latency exceeded 2x the low-load reference; loaded answer failed its quality check
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	not run for the 100K recipe
short prompt / long answer	short prompt -> 1,024-word target / 2,048-token ceiling do-sixhour-aimd-20260824-r1	healthy at 0.2 RPS once; not repeated	failed at 0.4 RPS: p95 time to first token exceeded 2x the low-load reference; low-load reference answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	repeatedly passed at 0.0666667 RPS; degraded by 0.2 RPS	failed at 0.1 RPS: p95 time to first token exceeded 2x the low-load reference; low-load reference answer failed its quality check

Capability	Transport	Functional
response format	observed supported	failed
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed transport degraded	not scored
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 182823; combined prompt + requested output: accepted and retrieval valid through 265289; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

kimi-k3

Model / API	kimi-k3 / chat_completions
Region / API version	not reported / v1
Context / max output	not documented / not documented tokens
Input / output price	\$2.85 / \$14.25 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	256-token prompt -> 128-token answer target do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	failed at 1 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
32K-token prompt / short answer	32,000-token prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	repeatedly passed through at least 0.8 RPS; ceiling not found	failed at 2.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
short prompt / long answer	short prompt -> 1,024-word target / 2,048-token ceiling do-sixhour-aimd-20260824-r1	healthy at 0.466667 RPS once; not repeated	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
seeded multi-workload mix	seeded four-way mix: 256->128, 100K/50K->128, 256->4,096, and 1,024->512 JSON do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	failed at 3.2 RPS: success rate fell below 99%; more than 1% of requests were rate limited

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	failed
parallel tool calls	observed supported	degraded
vision	observed supported	failed
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 158411; combined prompt + requested output: accepted and retrieval valid through 53633; output limit: realized generation observed; vision: observed functional; tool calling: observed functional; structured output: observed functional

mimo-v2.5-pro

Model / API	mimo-v2.5-pro / chat_completions
Region / API version	not reported / v1
Context / max output	262144 / 262144 tokens
Input / output price	\$0.40 / \$1.50 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	short exact-answer prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	repeatedly passed through at least 8 RPS; ceiling not found	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
32K-token prompt / short answer	32,000-token prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	healthy at 0.8 RPS once; not repeated	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
short prompt / long answer	256-token prompt -> 4,096-token answer target do-capacity-20260828-r2	repeatedly passed through at least 1 RPS; ceiling not found	failed at 0.4 RPS: p95 time to first token exceeded 2x the low-load reference; low-load reference answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	healthy at 0.302163 RPS once; not repeated	failed at 0.2 RPS: p95 time to first token exceeded 2x the low-load reference; low-load reference answer failed its quality check

Capability	Transport	Functional
response format	observed supported	passed
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed transport degraded	not scored
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 245518; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

minimax-m2.5

Model / API	minimax-m2.5 / chat_completions
Region / API version	not reported / v1
Context / max output	65536 / 65536 tokens
Input / output price	\$0.30 / \$1.20 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	256-token prompt -> 128-token answer target do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	failed at 16 RPS: success rate fell below 99%; more than 1% of requests were rate limited
32K-token prompt / short answer	32,000-token prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	repeatedly passed at 0.466667 RPS; degraded by 0.8 RPS	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
short prompt / long answer	short prompt -> 1,024-word target / 2,048-token ceiling do-sixhour-aimd-20260824-r1	repeatedly passed through at least 0.466667 RPS; ceiling not found	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	repeatedly passed through at least 1 RPS; ceiling not found	failed at 0.2 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	degraded
parallel tool calls	observed supported	degraded
vision	observed transport degraded	not scored
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 166688; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

nemotron-3-ultra-550b

Model / API	nemotron-3-ultra-550b / chat_completions
Region / API version	not reported / v1
Context / max output	131072 / 131072 tokens
Input / output price	\$0.90 / \$1.70 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	short exact-answer prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	healthy at 0.99393 RPS once; not repeated	failed at 1 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	not run for the 100K recipe
short prompt / long answer	256-token prompt -> 4,096-token answer target do-capacity-20260828-r2	repeatedly passed through at least 1 RPS; ceiling not found	failed at 1 RPS: p95 end-to-end latency exceeded 2x the low-load reference; low-load reference answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	repeatedly passed through at least 0.2 RPS; ceiling not found	failed at 0.4 RPS: p95 end-to-end latency exceeded 2x the low-load reference; low-load reference answer failed its quality check

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	degraded
vision	observed transport degraded	not scored
automatic prompt cache	documented supported	degraded
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 76667; combined prompt + requested output: accepted and retrieval valid through 126176; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

nvidia-nemotron-3-super-120b

Model / API	nvidia-nemotron-3-super-120b / chat_completions
Region / API version	not reported / v1
Context / max output	1000000 / 32768 tokens
Input / output price	\$0.30 / \$0.65 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	256-token prompt -> 128-token answer target do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	could not start: a reliable low-load baseline could not be established
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	not run for the 100K recipe
short prompt / long answer	short prompt -> 1,024-word target / 2,048-token ceiling do-sixhour-aimd-20260824-r1	healthy at 0.024582 RPS once; not repeated	failed at 0.1 RPS: p95 end-to-end latency exceeded 2x the low-load reference; low-load reference answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	no healthy result at the lowest tested rate (0.0666667 RPS)	could not start: a reliable low-load baseline could not be established

Capability	Transport	Functional
response format	not measured	not measured
tools	observed transport degraded	not scored
parallel tool calls	observed transport degraded	not scored
vision	observed transport degraded	not scored
automatic prompt cache	documented supported	degraded
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 76667; combined prompt + requested output: untested or inconclusive; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: inconclusive; structured output: observed functional

openai-gpt-oss-120b

Model / API	openai-gpt-oss-120b / chat_completions
Region / API version	not reported / v1
Context / max output	128000 / 4096 tokens
Input / output price	\$0.10 / \$0.70 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	short exact-answer prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	repeatedly passed through at least 8 RPS; ceiling not found	failed at 32 RPS: p95 time to first token exceeded 2x the low-load reference; p95 end-to-end latency exceeded 2x the low-load reference
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	not run for the 100K recipe
short prompt / long answer	256-token prompt -> 4,096-token answer target do-capacity-20260828-r2	test ran, but no defensible numeric rate was established	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	repeatedly passed at 0.0666667 RPS; degraded by 0.2 RPS	failed at 0.4 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	degraded
vision	observed transport degraded	not scored
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 66717; combined prompt + requested output: accepted and retrieval valid through 117823; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

qwen3.5-397b-a17b

Model / API	qwen3.5-397b-a17b / chat_completions
Region / API version	not reported / v1
Context / max output	131072 / 131072 tokens
Input / output price	\$0.55 / \$3.50 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	short exact-answer prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	no healthy result at the lowest tested rate (1 RPS)	failed at 16 RPS: success rate fell below 99%; more than 1% of requests were rate limited
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	not run for the 100K recipe
short prompt / long answer	256-token prompt -> 4,096-token answer target do-capacity-20260828-r2	no healthy result at the lowest tested rate (0.25 RPS)	failed at 0.2 RPS: low-load reference answer failed its quality check; loaded answer failed its quality check
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	no healthy result at the lowest tested rate (0.0666667 RPS)	could not start: a reliable low-load baseline could not be established

Capability	Transport	Functional
response format	not measured	not measured
tools	observed transport degraded	not scored
parallel tool calls	observed transport degraded	not scored
vision	observed supported	failed
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: nonmonotonic context outcomes inconclusive; combined prompt + requested output: accepted and retrieval valid through 70999; output limit: realized generation observed; vision: inconclusive; tool calling: inconclusive; structured output: observed functional

qwen3.8-max

Model / API	qwen3.8-max / chat_completions
Region / API version	not reported / v1
Context / max output	1000000 / 262144 tokens
Input / output price	\$2.00 / \$6.00 per million tokens

Workload	Exact recipe and source	Adaptive-load result	120 s fixed-rate result
short prompt / short answer	short exact-answer prompt -> 64-token answer ceiling do-sixhour-aimd-20260824-r1	healthy at 2.46713 RPS once; not repeated	PASS at 1 RPS
100K-token prompt / short answer	100,000-token prompt -> 128-token answer target (50,000-token prompt for Minimax M2.5) do-capacity-20260828-r2	repeatedly passed through at least 0.5 RPS; ceiling not found	not run for the 100K recipe
short prompt / long answer	short prompt -> 1,024-word target / 2,048-token ceiling do-sixhour-aimd-20260824-r1	healthy at 0.913063 RPS once; not repeated	failed at 0.4 RPS: p95 time to first token exceeded 2x the low-load reference; p95 end-to-end latency exceeded 2x the low-load reference
seeded multi-workload mix	seeded five-way mix: short exact answer, 4,096-token context, 512-word answer, JSON, and tool call; 1,024-token ceiling do-sixhour-aimd-20260824-r1	healthy at 2.00595 RPS once; not repeated	failed at 0.1 RPS: low-load reference answer failed its quality check; p95 end-to-end latency exceeded 2x the low-load reference

Capability	Transport	Functional
response format	observed supported	passed
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed supported	passed
automatic prompt cache	documented supported	passed
batch open models	documented unavailable	documented unavailable

Observed boundaries: prompt context window: accepted and retrieval valid through 1004307; combined prompt + requested output: accepted and retrieval valid through 1002293; output limit: realized generation observed; vision: observed functional; tool calling: observed functional; structured output: observed functional