

DigitalOcean hosted inference

Technical evidence atlas

12 exact hosted endpoints. Capacity source: do-sixhour-aimd-20260824-r1. Sustained-load source: do-direct-soak-20260823-r1.

The AIMD pages show tested operational bounds, not fitted theoretical maxima. The soak pages report four independent 30-second blocks at one candidate rate. Missing values remain blank; singleton boundary/capability probes do not receive invented confidence intervals.


Engineering decision map

This inventory reports evidence states for each exact hosted endpoint. It is not a single global ranking: capacity and capability must be read for the workload and feature the application actually uses.

Exact endpoint	Context	Confirmed AIMD	Complete soaks	Capabilities passed
arcee-trinity-large-thinking	128000	2/4	4/4	1/17
deepseek-v4-flash-0731	1048576	4/4	4/4	11/17
gemma-4-31B-it	256000	2/4	4/4	12/17
glm-5.2	262144	1/4	4/4	7/24
kimi-k3	not documented	2/4	4/4	1/17
mimo-v2.5-pro	262144	1/4	4/4	9/24
minimax-m2.5	65536	4/4	4/4	0/17
nemotron-3-ultra-550b	131072	2/4	4/4	9/17
nvidia-nemotron-3-super-120b	1000000	1/4	2/4	7/17
openai-gpt-oss-120b	128000	3/4	4/4	5/17
qwen3.5-397b-a17b	131072	0/4	3/4	0/17
qwen3.8-max	1000000	0/4	4/4	17/24


DigitalOcean AIMD capacity — short input / short output

confirmed healthy lower bound exploratory healthy observation unhealthy upper bound




DigitalOcean AIMD capacity — 32K input / short output

● confirmed healthy lower bound □ exploratory healthy observation | unhealthy upper bound




DigitalOcean AIMD capacity — short input / long output


confirmed healthy lower bound exploratory healthy observation unhealthy upper bound



DigitalOcean AIMD capacity — heterogeneous production mix




DigitalOcean two-minute soak — short input / short output




Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean two-minute soak — short input / short output




Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean two-minute soak — 32K input / short output




Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean two-minute soak — 32K input / short output




Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean two-minute soak — short input / long output




Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean two-minute soak — short input / long output




Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean two-minute soak — heterogeneous production mix




Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean two-minute soak — heterogeneous production mix



Whiskers are 95% Student-t intervals across four 30-second blocks. Missing estimates are labelled, never plotted as zero.

DigitalOcean functional capability evidence



arcee-trinity-large-thinking

Model / API	arcee-trinity-large-thinking / chat_completions
Region / API version	not reported / v1
Context / max output	128000 / 32000 tokens
Input / output price	\$0.25 / \$0.90 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	confirmed right censored lower bound	16 RPS	1 RPS / complete
32K input / short output	confirmed right censored lower bound	0.8 RPS	0.4 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	1 RPS / complete
heterogeneous production mix	unconfirmed healthy observation only	not established	0.4 RPS / complete

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed transport degraded	not scored
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 174988; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

deepseek-v4-flash-0731

Model / API	deepseek-v4-flash-0731 / chat_completions
Region / API version	not reported / v1
Context / max output	1048576 / 1048576 tokens
Input / output price	\$0.080 / \$0.252 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	confirmed right censored lower bound	32 RPS	1 RPS / complete
32K input / short output	confirmed right censored lower bound	0.8 RPS	0.4 RPS / complete
short input / long output	confirmed right censored lower bound	0.467 RPS	1.6 RPS / complete
heterogeneous production mix	confirmed bracketed interval	0.0667 RPS	1.6 RPS / complete

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 2722; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

gemma-4-31B-it

Model / API	gemma-4-31B-it / chat_completions
Region / API version	not reported / v1
Context / max output	256000 / 8192 tokens
Input / output price	\$0.18 / \$0.50 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	unconfirmed healthy observation only	not established	1 RPS / complete
32K input / short output	confirmed bracketed interval	0.0667 RPS	0.4 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	0.4 RPS / complete
heterogeneous production mix	confirmed bracketed interval	0.0667 RPS	0.4 RPS / complete

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed supported	passed

Observed boundaries: prompt context window: accepted and retrieval valid through 251873; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed functional; tool calling: observed functional; structured output: observed functional

glm-5.2

Model / API	glm-5.2 / chat_completions
Region / API version	not reported / v1
Context / max output	262144 / 262144 tokens
Input / output price	\$0.70 / \$2.20 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	unconfirmed healthy observation only	not established	1 RPS / complete
32K input / short output	unconfirmed healthy observation only	not established	0.4 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	0.4 RPS / complete
heterogeneous production mix	confirmed bracketed interval	0.0667 RPS	0.1 RPS / complete

Capability	Transport	Functional
response format	observed supported	failed
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 182823; combined prompt + requested output: accepted and retrieval valid through 265289; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

kimi-k3

Model / API	kimi-k3 / chat_completions
Region / API version	not reported / v1
Context / max output	not documented / not documented tokens
Input / output price	\$2.85 / \$14.25 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	unconfirmed healthy observation only	not established	1 RPS / complete
32K input / short output	confirmed right censored lower bound	0.8 RPS	2.4 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	0.4 RPS / complete
heterogeneous production mix	confirmed bracketed interval	0.0667 RPS	3.2 RPS / complete

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	failed
parallel tool calls	observed supported	degraded
vision	observed supported	failed

Observed boundaries: prompt context window: accepted and retrieval valid through 158411; combined prompt + requested output: accepted and retrieval valid through 53633; output limit: realized generation observed; vision: observed functional; tool calling: observed functional; structured output: observed functional

mimo-v2.5-pro

Model / API	mimo-v2.5-pro / chat_completions
Region / API version	not reported / v1
Context / max output	262144 / 262144 tokens
Input / output price	\$0.40 / \$1.50 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	confirmed right censored lower bound	8 RPS	0.4 RPS / complete
32K input / short output	unconfirmed healthy observation only	not established	0.4 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	0.4 RPS / complete
heterogeneous production mix	unconfirmed healthy observation only	not established	0.2 RPS / complete

Capability	Transport	Functional
response format	observed supported	passed
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 245518; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

minimax-m2.5

Model / API	minimax-m2.5 / chat_completions
Region / API version	not reported / v1
Context / max output	65536 / 65536 tokens
Input / output price	\$0.30 / \$1.20 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	confirmed bracketed interval	8 RPS	16 RPS / complete
32K input / short output	confirmed bracketed interval	0.467 RPS	0.4 RPS / complete
short input / long output	confirmed right censored lower bound	0.467 RPS	0.4 RPS / complete
heterogeneous production mix	confirmed right censored lower bound	1 RPS	0.2 RPS / complete

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	degraded
parallel tool calls	observed supported	degraded
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 166688; combined prompt + requested output: transport accepted but retrieval not verified; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

nemotron-3-ultra-550b

Model / API	nemotron-3-ultra-550b / chat_completions
Region / API version	not reported / v1
Context / max output	131072 / 131072 tokens
Input / output price	\$0.90 / \$1.70 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	unconfirmed healthy observation only	not established	1 RPS / complete
32K input / short output	confirmed bracketed interval	0.2 RPS	0.8 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	1 RPS / complete
heterogeneous production mix	confirmed right censored lower bound	0.2 RPS	0.4 RPS / complete

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	degraded
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 76667; combined prompt + requested output: accepted and retrieval valid through 126176; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

nvidia-nemotron-3-super-120b

Model / API	nvidia-nemotron-3-super-120b / chat_completions
Region / API version	not reported / v1
Context / max output	1000000 / 32768 tokens
Input / output price	\$0.30 / \$0.65 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	confirmed bracketed interval	1 RPS	not measured / baseline transport gate failed
32K input / short output	censored no valid healthy epoch	not established	0.4 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	0.1 RPS / complete
heterogeneous production mix	censored no valid healthy epoch	not established	not measured / baseline transport gate failed

Capability	Transport	Functional
response format	not measured	not measured
tools	observed transport degraded	not scored
parallel tool calls	observed transport degraded	not scored
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 76667; combined prompt + requested output: untested or inconclusive; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: inconclusive; structured output: observed functional

openai-gpt-oss-120b

Model / API	openai-gpt-oss-120b / chat_completions
Region / API version	not reported / v1
Context / max output	128000 / 4096 tokens
Input / output price	\$0.10 / \$0.70 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	confirmed right censored lower bound	8 RPS	32 RPS / complete
32K input / short output	confirmed right censored lower bound	0.8 RPS	2.4 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	0.4 RPS / complete
heterogeneous production mix	confirmed bracketed interval	0.0667 RPS	0.4 RPS / complete

Capability	Transport	Functional
response format	not measured	not measured
tools	observed supported	passed
parallel tool calls	observed supported	degraded
vision	observed transport degraded	not scored

Observed boundaries: prompt context window: accepted and retrieval valid through 66717; combined prompt + requested output: accepted and retrieval valid through 117823; output limit: realized generation observed; vision: observed rejected or unsupported; tool calling: observed functional; structured output: observed functional

qwen3.5-397b-a17b

Model / API	qwen3.5-397b-a17b / chat_completions
Region / API version	not reported / v1
Context / max output	131072 / 131072 tokens
Input / output price	\$0.55 / \$3.50 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	censored no valid healthy epoch	not established	16 RPS / complete
32K input / short output	censored no valid healthy epoch	not established	0.4 RPS / complete
short input / long output	censored no valid healthy epoch	not established	0.2 RPS / complete
heterogeneous production mix	censored no valid healthy epoch	not established	not measured / baseline transport gate failed

Capability	Transport	Functional
response format	not measured	not measured
tools	observed transport degraded	not scored
parallel tool calls	observed transport degraded	not scored
vision	observed supported	failed

Observed boundaries: prompt context window: nonmonotonic context outcomes inconclusive; combined prompt + requested output: accepted and retrieval valid through 70999; output limit: realized generation observed; vision: inconclusive; tool calling: inconclusive; structured output: observed functional

qwen3.8-max

Model / API	qwen3.8-max / chat_completions
Region / API version	not reported / v1
Context / max output	1000000 / 262144 tokens
Input / output price	\$2.00 / \$6.00 per million tokens

Workload	AIMD evidence	Healthy bound	Soak rate / state
short input / short output	unconfirmed healthy observation only	not established	1 RPS / complete
32K input / short output	unconfirmed healthy observation only	not established	0.2 RPS / complete
short input / long output	unconfirmed healthy observation only	not established	0.4 RPS / complete
heterogeneous production mix	unconfirmed healthy observation only	not established	0.1 RPS / complete

Capability	Transport	Functional
response format	observed supported	passed
tools	observed supported	passed
parallel tool calls	observed supported	passed
vision	observed supported	passed

Observed boundaries: prompt context window: accepted and retrieval valid through 1004307; combined prompt + requested output: accepted and retrieval valid through 1002293; output limit: realized generation observed; vision: observed functional; tool calling: observed functional; structured output: observed functional